

P5- ENRON SUBMISSION

FREE RESPONSE QUESTIONS

1. Summarize for us the goal of this project and how machine learning is useful in trying to accomplish it. As part of your answer, give some background on the dataset and how it can be used to answer the project question. Were there any outliers in the data when you got it, and how did you handle those? [relevant rubric items: "data exploration", "outlier investigation"]

Enron Corporation was an American energy company which was once ranked sixth-largest energy company in the world. At Enron's peak, its shares were worth \$90.75, but after the company declared bankruptcy on December 2, 2001, they plummeted to \$0.67 by January 2002 due to widespread corporate fraud.

To this day, many wonder how such a powerful business disintegrated almost overnight and how it managed to fool the regulators with fake, off-the-books corporations for so long. As part of this project, I would try investigating the Enron fraud using machine learning skills I learned to try building a person of interest identifier based on financial and email data made public as a result of the Enron scandal.

The dataset contained:

Number of data points: 146
Number of person of interest: 18
Number of non-person of Interest: 128
Number of features per person: 21

Missing values :

Salary	51
to_messages	60
deferral_payments	107
total_payments	21
long_term_incentive	80
loan_advances	142
Bonus	64
restricted_stock	36
restricted_stock_deferred	128
total_stock_value	20
Expenses	51
from_poi_to_this_person	60
exercised_stock_options	44
from_messages	60
Other	53
from_this_person_to_poi	60

Poi	0
deferred_income	97
shared_receipt_with_poi	60
email_address	35
director_fees	129

I identified the below 3 outliers in the data set through exploratory data analysis:

- **TOTAL**: This outlier was the sum total of all the payments and stock values carried out by people in the Enron financial data set. So it didn't hold any further significance in my analysis, so I removed it.
 - **THE TRAVEL AGENCY IN THE PARK**: This outlier doesn't seem to correspond to a real person with all the payments and stock value data point having 0 values barring other expenses and total payments so it this data point has also been removed from the data set.
 - **LOCKHART EUGENE E**: This person has all the payments and stock value data point values as 0 so I removed it from the data set.
2. What features did you end up using in your POI identifier, and what selection process did you use to pick them? Did you have to do any scaling? Why or why not? As part of the assignment, you should attempt to engineer your own feature that does not come ready-made in the dataset -- explain what feature you tried to make, and the rationale behind it. (You do not necessarily have to use it in the final analysis, only engineer and test it.) In your feature selection step, if you used an algorithm like a decision tree, please also give the feature importances of the features that you use, and if you used an automated feature selection function like SelectKBest, please report the feature scores and reasons for your choice of parameter values. [relevant rubric items: "create new features", "properly scale features", "intelligently select feature"]

I believe for people to be involved in the fraud, they must have had some financial communication with POI. So to identify features like this which could help in my further analysis, I created the following 3 new features on the basis of e-mails exchanged between non POI and POI:

- **fraction_from_poi**: proportion of messages received from POI
- **fraction_to_poi**: proportion of messages sent to POI
- **messages_with_poi**: sum total of messages exchanged with POI

I used sklearn's SelectKBest for feature selection to select top k best features with k from 5 to 10. I chose the default value of k=10 for my feature selection, the scores come out to be:

k=10:
exercised_stock_options: 24.815080
total_stock_value: 24.182899
bonus: 20.792252
salary: 18.289684
fraction_to_poi: 16.409713
deferred_income: 11.458477
long_term_incentive: 9.922186

restricted_stock: 9.212811
total_payments: 8.772778
messages_with_poi: 8.616648

To scale the features, I used MinMaxScaler from sklearn to normalize the features as financial/email data numbers vary in high magnitude.

3. What algorithm did you end up using? What other one(s) did you try? How did model performance differ between algorithms? [relevant rubric item: "pick an algorithm"]

I tried algorithms like **Nearest Neighbors, SVM(Linear/Rbf), Decision Tree, Random Fores, AdaBoost, Naive Bayes, Extra Trees** for my classifiers with their default values. Following are the result of the algorithms:

Algorithm	Accuracy	Precision	Recall
K Nearest Neighbors	0. 883721	0.0	0.0
Linear SVM	0. 883721	0.0	0.0
RBF SVM	0. 883721	0.0	0.0
Decision Tree	0. 883721	0. 500000	0. 400000
Random Forest	0. 883721	0. 500000	0. 400000
AdaBoost	0. 860465	0. 400000	0. 400000
Naive Bayes	0. 883721	0.500000	0. 600000
Extra Trees	0. 837209	0. 250000	0. 200000

4. What does it mean to tune the parameters of an algorithm, and what can happen if you don't do this well? How did you tune the parameters of your particular algorithm? (Some algorithms do not have parameters that you need to tune -- if this is the case for the one you picked, identify and briefly explain how you would have done it for the model that was not your final choice or a different model that does utilize parameter tuning, e.g. a decision tree classifier). [relevant rubric item: "tune the algorithm"]
5. What is validation, and what's a classic mistake you can make if you do it wrong? How did you validate your analysis? [relevant rubric item: "validation strategy"]
6. Give at least 2 evaluation metrics and your average performance for each of them. Explain an interpretation of your metrics that says something human-understandable about your algorithm's performance. [relevant rubric item: "usage of evaluation metrics"]